

WeCanDoNow HandsOn Lab - Kafka Streams API

by Saravanan Sundaramoorthy

The Streams API allows transforming streams of data from input topics to output topics.

```
sudo apt install maven -y

mvn archetype:generate -DgroupId=com.wecandonow -DartifactId=WeCanDoNow-
Kafka-Streams -DarchetypeArtifactId=maven-archetype-quickstart
cd WeCanDoNow-Kafka-Streams/
```

Add Kafka dependencies in `pom.xml`

```
cd WeCanDoNow-Kafka-Streams

vim pom.xml
```

```
<project xmlns="http://maven.apache.org/POM/4.0.0"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
xsi:schemaLocation="http://maven.apache.org/POM/4.0.0
http://maven.apache.org/maven-v4_0_0.xsd">
  <modelVersion>4.0.0</modelVersion>
  <groupId>com.wecandonow</groupId>
  <artifactId>Kafka-Producer</artifactId>
  <packaging>jar</packaging>
  <version>1.0-SNAPSHOT</version>
  <name>Kafka-Producer</name>
  <url>http://maven.apache.org</url>

  <build>
    <plugins>
      <plugin>
        <groupId>org.apache.maven.plugins</groupId>
        <artifactId>maven-compiler-plugin</artifactId>
        <version>3.7.0</version>
        <configuration>
          <source>1.8</source>
          <target>1.8</target>
        </configuration>
      </plugin>
      <plugin>
        <groupId>org.apache.maven.plugins</groupId>
        <artifactId>maven-assembly-plugin</artifactId>
        <version>2.4.1</version>
```

```
        <configuration>
          <archive>
            <manifest>
              <mainClass>com.wecandonow.App</mainClass>
            </manifest>
          </archive>
          <descriptorRefs>
            <descriptorRef>jar-with-dependencies</descriptorRef>
          </descriptorRefs>
          <skipAssembly>>false</skipAssembly>
        </configuration>
      <executions>
        <execution>
          <id>package</id>
          <phase>package</phase>
          <goals>
            <goal>single</goal>
          </goals>
        </execution>
      </executions>
    </plugin>
  </plugins>
</build>

<dependencies>
  <dependency>
    <groupId>junit</groupId>
    <artifactId>junit</artifactId>
    <version>3.8.1</version>
    <scope>test</scope>
  </dependency>

  <dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka_2.12</artifactId>
    <version>1.0.1</version>
  </dependency>

  <dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>
    <version>1.0.1</version>
  </dependency>

  <dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-streams</artifactId>
    <version>1.0.1</version>
  </dependency>
</dependencies>
</project>
```

Streams.java

```
vim src/main/java/com/wecandonow/Streams.java
```

Reads from input topic and sends message length to output topic

```
package com.wecandonow;

import org.apache.kafka.common.serialization.Serdes;
import org.apache.kafka.streams.KafkaStreams;
import org.apache.kafka.streams.StreamsBuilder;
import org.apache.kafka.streams.StreamsConfig;

import java.util.HashMap;
import java.util.Map;

public class Streams
{
    public static void main(String[] args) {
        if (args.length != 2) {
            System.out.println("Usage: <input_topic> <output_topic>");
            return;
        }

        Map<String, Object> props = new HashMap<>();
        props.put(StreamsConfig.APPLICATION_ID_CONFIG, "my-stream-
processing-application");
        props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG,
"localhost:9092");
        props.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG,
Serdes.String().getClass());
        props.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG,
Serdes.String().getClass());
        StreamsConfig config = new StreamsConfig(props);

        StreamsBuilder builder = new StreamsBuilder();
        builder.<String, String>stream(args[0]).mapValues(value ->
value.length() + "").to(args[1]);

        KafkaStreams streams = new KafkaStreams(builder.build(),
config);
        streams.start();
    }
}
```

Create input and output topics

```
kafka-topics.sh --delete --zookeeper localhost:2181 --topic input-topic
kafka-topics.sh --delete --zookeeper localhost:2181 --topic output-topic

kafka-topics.sh --create --zookeeper localhost:2181 --topic input-topic --
partitions 1 --replication-factor 1
kafka-topics.sh --create --zookeeper localhost:2181 --topic output-topic -
-partitions 1 --replication-factor 1
```

Compile and Run the App

```
mvn compile
mvn exec:java -Dexec.mainClass="com.wecandonow.Streams" -
Dexec.args="input-topic output-topic"
```

Output

```
kafka-console-producer.sh --broker-list localhost:9092 --topic input-topic
>hello
>world
>wecandonow
```

```
kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic
output-topic --from-beginning
5
5
7
```